

# Bike Sharing MLOps Project - Complete Documentation

## Table of Contents


1. [Project Overview](#project-overview)
2. [Architecture](#architecture)
3. [File-by-File Analysis](#file-by-file-analysis)
4. [Usage Guide](#usage-guide)
5. [Deployment Instructions](#deployment-instructions)

## Project Overview

This is a complete MLOps (Machine Learning Operations) project for predicting bike-sharing demand. It implements:

- **ML Pipeline**: Data ingestion, validation, training, prediction
- **Orchestration**: Apache Airflow for workflow automation
- **Model Registry**: MLflow for experiment tracking and model versioning
- **Monitoring**: Evidently AI for data drift detection, Prometheus & Grafana for metrics
- **API**: FastAPI for serving predictions
- **Frontend**: Streamlit for user interface
- **Infrastructure**: Docker containers, PostgreSQL database, LocalStack (AWS S3 simulation)
- **CI/CD**: GitHub Actions for automated testing and deployment
- **Data Versioning**: DVC (Data Version Control)

## Architecture





## File-by-File Analysis

### 1. DOCKER CONFIGURATION FILES

#### *\*\*docker-compose.yml\*\**

**Purpose:** Orchestrates all services in the MLOps pipeline using Docker containers.

#### Line-by-Line Explanation:

```
version: '3.8'
```

- Specifies Docker Compose file format version
- **Why**: Version 3.8 supports all modern Docker features
- **Usage**: Required at the top of every docker-compose file

```
services:
  postgres:
    image: postgres:13
```

- Defines PostgreSQL database service
- **Why**: Stores Airflow metadata and prediction logs
- **Usage**: Database for tracking all predictions and workflow states

```
  container_name: postgres_db
  networks: [bike-mlops-net]
  ports: ["5432:5432"]
```

- `container\_name`: Custom name for easy reference
- `networks`: Connects to shared network for inter-container communication
- `ports`: Maps container port 5432 to host port 5432
- **Why**: Allows external tools (like pgAdmin) to connect
- **Usage**: Access database at `localhost:5432`

```

environment:
  - POSTGRES_USER=airflow
  - POSTGRES_PASSWORD=airflow
  - POSTGRES_DB=airflow

```

- Sets database credentials and default database name
- \*\*Why\*\*: Airflow requires these specific credentials
- \*\*Usage\*\*: Use these credentials in connection strings

```

healthcheck:
  test: [ "CMD-SHELL", "pg_isready -U airflow" ]
  interval: 5s
  timeout: 5s
  retries: 5

```

- Checks if PostgreSQL is ready to accept connections
- \*\*Why\*\*: Prevents other services from starting before database is ready
- \*\*Usage\*\*: Other services use `depends\_on: postgres: {condition: service\_healthy}`

```

localstack:
  image: localstack/localstack:latest
  container_name: localstack
  ports: [ "4566:4566" ]

```

- LocalStack simulates AWS services locally
- \*\*Why\*\*: Test S3 storage without AWS costs
- \*\*Usage\*\*: Upload monitoring reports to simulated S3 bucket

```

environment:
  - SERVICES=s3

```

- Only enables S3 service (not EC2, Lambda, etc.)
- \*\*Why\*\*: Reduces resource usage
- \*\*Usage\*\*: Access at `http://localhost:4566`

```

volumes:
  - "./localstack_data:/var/lib/localstack"

```

- Persists LocalStack data on host machine
- \*\*Why\*\*: Data survives container restarts
- \*\*Usage\*\*: S3 buckets and objects stored in `./localstack\_data`

```

airflow-init:
  image: apache/airflow:2.7.1
  container_name: airflow_init
  depends_on:
    postgres: {condition: service_healthy}

```

- One-time initialization container for Airflow
- \*\*Why\*\*: Creates database schema and admin user
- \*\*Usage\*\*: Runs once, then exits

```

command: >
  bash -c "airflow db init &&
  airflow users create --username admin --firstname admin --lastname admin --role Admin --email admin@e"

```

- Initializes Airflow database and creates admin user
- \*\*Why\*\*: Required before Airflow can start
- \*\*Usage\*\*: Login to Airflow UI with `admin/admin`

```
mlflow:
  image: ghcr.io/mlflow/mlflow:v2.8.1
  container_name: mlflow_server
  ports: [ "5000:5000" ]
```

- MLflow tracking server for experiment management
- \*\*Why\*\*: Tracks model versions, parameters, metrics
- \*\*Usage\*\*: Access UI at `http://localhost:5000`

```
volumes:
  - ./mlflow_data:/mlflow_db
  - ./mlruns:/mlflow/mlruns
```

- Persists MLflow data
- \*\*Why\*\*: Keeps experiment history across restarts
- \*\*Usage\*\*: All model artifacts stored in `./mlruns`

```
command: mlflow server --host 0.0.0.0 --port 5000 --backend-store-uri sqlite:///mlflow_db/mlflow.db --
```

- Starts MLflow server with SQLite backend
- \*\*Why\*\*: Lightweight database for tracking
- \*\*Usage\*\*: Python code connects to `http://mlflow\_server:5000`

```
webserver:
  build:
    context: .
    dockerfile: Dockerfile.airflow
```

- Builds custom Airflow image with additional dependencies
- \*\*Why\*\*: Needs pandas, evidently, mlflow not in base image
- \*\*Usage\*\*: Airflow UI accessible at `http://localhost:8081`

```
command: webserver
extra_hosts: [ "host.docker.internal:host-gateway" ]
```

- `command`: Starts Airflow web interface
- `extra\_hosts`: Allows container to access host machine
- \*\*Why\*\*: API runs on Windows host, containers need to reach it
- \*\*Usage\*\*: Use `http://host.docker.internal:9999` in container code

```
depends_on:
  airflow-init: {condition: service_completed_successfully}
  mlflow: {condition: service_started}
```

- Ensures initialization completes before starting
- \*\*Why\*\*: Prevents startup errors

- **\*\*Usage\*\*:** Automatic dependency management

```
environment: &airflow-common-env
  - AIRFLOW__DATABASE__SQLALCHEMY_CONN=postgresql+psycopg2://airflow:airflow@postgres/airflow
  - AIRFLOW__CORE__LOAD_EXAMPLES=False
  - AIRFLOW__WEBSERVER__SECRET_KEY=my_secret_key_123
```

- `&airflow-common-env;`: YAML anchor for reuse
- Database connection string
- Disables example DAGs
- Secret key for session management
- **\*\*Why\*\*:** Configuration shared between webserver and scheduler
- **\*\*Usage\*\*:** Referenced as `\*airflow-common-env` in scheduler

```
ports: [ "8081:8080" ]
```

- Maps Airflow UI to port 8081 on host
- **\*\*Why\*\*:** Port 8080 might be used by other services
- **\*\*Usage\*\*:** Access at `http://localhost:8081`

```
volumes: &airflow-common-volumes
  - ./dags:/opt/airflow/dags
  - ./src:/opt/airflow/src
  - ./data:/opt/airflow/data
  - ./models:/opt/airflow/models
  - ./mlruns:/mlruns
```

- Mounts local directories into container
- **\*\*Why\*\*:** Code changes reflect immediately without rebuild
- **\*\*Usage\*\*:** Edit files locally, Airflow sees changes instantly

```
scheduler:
  build:
    context: .
    dockerfile: Dockerfile.airflow
  container_name: airflow_scheduler
  command: scheduler
  user: root
```

- Airflow scheduler triggers DAG runs
- **\*\*Why\*\*:** Monitors DAGs and executes tasks on schedule
- **\*\*Usage\*\*:** Runs in background, no direct interaction
- `user: root`: Needed for file permissions

```
bike_sharing_frontend:
  build:
    context: .
    dockerfile: Dockerfile
  container_name: streamlit_app
  ports: [ "8501:8501" ]
```

- Streamlit web application
- **\*\*Why\*\*:** User-friendly interface for predictions
- **\*\*Usage\*\*:** Access at `http://localhost:8501`

```
environment:
  - API_URL=http://host.docker.internal:9999
```

- Tells Streamlit where to find the API
- **Why**: API runs on host, not in container
- **Usage**: Streamlit makes requests to this URL

```
volumes:
  - ./src:/opt/airflow/src
command: ["streamlit", "run", "/opt/airflow/src/app.py", "--server.port=8501", "--server.address=0.0.0.0"]
```

- Mounts source code
- Starts Streamlit app
- **Why**: Allows code updates without rebuild
- **Usage**: Edit `app.py`, refresh browser to see changes

```
prometheus:
  image: prom/prometheus:latest
  container_name: prometheus
  volumes:
    - ./prometheus.yml:/etc/prometheus/prometheus.yml
  ports: [ "9090:9090" ]
```

- Prometheus metrics collection system
- **Why**: Monitors API performance and request counts
- **Usage**: Access at `http://localhost:9090`

```
grafana:
  image: grafana/grafana:latest
  container_name: grafana
  ports: [ "3000:3000" ]
  depends_on: [prometheus]
```

- Grafana visualization dashboard
- **Why**: Beautiful charts for Prometheus metrics
- **Usage**: Access at `http://localhost:3000` (default: admin/admin)

```
networks:
  bike-mlops-net:
    driver: bridge
```

- Creates isolated network for all services
- **Why**: Containers can communicate by service name
- **Usage**: Use `postgres`, `mlflow\_server`, etc. as hostnames

**When to Use:** Run `docker-compose up -d` to start all services

## **\*\*Dockerfile\*\***

**Purpose:** Builds the Streamlit frontend container.

```
FROM python:3.9-slim
```

- Base image with Python 3.9
- **Why**: Lightweight image reduces build time
- **Usage**: Foundation for all Python dependencies

```
WORKDIR /app
```

- Sets working directory inside container
- **Why**: All subsequent commands run from this directory
- **Usage**: Files copied to `/app`

```
COPY src/requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt
```

- Copies requirements file first
- Installs Python packages
- **Why**: Docker caches this layer, speeds up rebuilds
- **Usage**: Only rebuilds if requirements.txt changes

```
COPY . .
```

- Copies all project files
- **Why**: Needed for Streamlit app to run
- **Usage**: Includes src/, data/, models/

```
EXPOSE 8501
```

- Documents that container listens on port 8501
- **Why**: Informational, doesn't actually open port
- **Usage**: Reminder for port mapping in docker-compose

```
CMD [ "streamlit", "run", "src/app.py", "--server.port=8501", "--server.address=0.0.0.0" ]
```

- Default command when container starts
- **Why**: Launches Streamlit application
- **Usage**: Overridden by docker-compose command

**When to Use:** Automatically built by docker-compose

### ***\*\*Dockerfile.airflow\*\****

**Purpose:** Builds custom Airflow image with ML dependencies.

```
FROM apache/airflow:2.7.1-python3.9
```

- Official Airflow base image
- **Why**: Pre-configured with Airflow dependencies
- **Usage**: Adds ML libraries on top

```
USER root
RUN apt-get update && apt-get install -y gcc python3-dev
```

- Switches to root user
- Installs system packages
- **Why**: gcc needed to compile some Python packages
- **Usage**: Required for psycopg2, pandas compilation

```
USER airflow
```

- Switches back to airflow user
- **Why**: Security best practice (don't run as root)
- **Usage**: All subsequent commands run as airflow user

```
RUN pip install --no-cache-dir --default-timeout=1000 \
    pandas \
    pydantic==1.10.13 \
    evidently==0.4.15 \
    sqlalchemy \
    psycopg2-binary \
    mlflow
```

- Installs ML and data science libraries
- `pydantic==1.10.13`: Specific version to avoid conflicts
- **Why**: Airflow base image doesn't include these
- **Usage**: Available in all Airflow tasks

**When to Use:** Automatically built by docker-compose

### ***\*\*.dockerignore\*\****

**Purpose:** Excludes files from Docker build context.

```
logs/
mlruns/
mlflow_data/
venv/
.git/
*.pyc
__pycache__/
```

- Lists files/folders to ignore during `docker build`
- **Why**: Reduces build context size, speeds up builds
- **Usage**: Prevents unnecessary files in image

**When to Use:** Automatically used during Docker builds

## **2. AIRFLOW DAG FILES**

## **\*\*dags/bike\_sharing\_dag.py\*\***

**Purpose:** Main ML pipeline orchestration - data ingestion, validation, training, and prediction.

```
from airflow import DAG
from airflow.operators.python import PythonOperator
from airflow.operators.bash import BashOperator
from airflow.utils.dates import days_ago
from datetime import datetime, timedelta
import sys
import os
import pandas as pd
import requests
```

- Imports Airflow components and utilities
- **Why**: Needed to define workflows and tasks
- **Usage**: Standard imports for any Airflow DAG

```
def call_live_api_predict():
    API_URL = "http://host.docker.internal:9999/predict"
```

- Function to test the live prediction API
- `host.docker.internal`: Special DNS name to reach host machine from Docker
- **Why**: API runs on Windows host, Airflow runs in container
- **Usage**: Called as Airflow task to verify API is working

```
params = {
    "season": 1, "mnth": 6, "hr": 10,
    "holiday": 0, "weekday": 3, "workingday": 1,
    "weathersit": 1, "temp": 0.5, "atemp": 0.5,
    "hum": 0.5, "windspeed": 0.1
}
```

- Sample input data for prediction
- **Why**: Tests API with realistic values
- **Usage**: Modify these to test different scenarios

```
response = requests.get(API_URL, params=params, timeout=60)
```

- Makes HTTP GET request to API
- `timeout=60`: Waits up to 60 seconds
- **Why**: Prevents task from hanging indefinitely
- **Usage**: Returns prediction result

```
if response.status_code == 200:
    print(f"■ API SUCCESS! Prediction Saved: {response.json()}")
else:
    print(f"■ API Error: Status {response.status_code}")
    raise Exception(f"API failed with status {response.status_code}")
```

- Checks if API call succeeded
- **Why**: Raises exception to mark Airflow task as failed
- **Usage**: Red task in Airflow UI indicates API problem

```
def run_inference_test():
    sys.path.append('/opt/airflow/src')
    from predict import make_prediction
```

- Internal prediction test (doesn't use API)
- **Why**: Verifies model works inside Airflow container
- **Usage**: Tests model directly without network calls

```
data_path = "/opt/airflow/data/processed/X_train.csv"
if os.path.exists(data_path):
    sample_df = pd.read_csv(data_path).sample(n=1)
    if 'yr' in sample_df.columns:
        sample_df = sample_df.drop(columns=['yr'])
```

- Loads one random sample from training data
- Removes 'yr' column if present
- **Why**: Model was trained without 'yr' column
- **Usage**: Ensures feature consistency

```
result = make_prediction(sample_df)
print(f"■ Internal Test Result: {result}")
```

- Makes prediction using loaded model
- **Why**: Validates model file is accessible and working
- **Usage**: Prints prediction to Airflow logs

```
default_args = {
    'owner': 'rama',
    'retries': 0,
    'retry_delay': timedelta(minutes=5),
}
```

- Default settings for all tasks in DAG
- `retries: 0`: Don't retry failed tasks
- **Why**: Easier to debug errors immediately
- **Usage**: Change to `retries: 3` for production

```
with DAG(
    'bike_sharing_final_pipeline_v4',
    default_args=default_args,
    description='Final Pipeline with host.docker.internal fix',
    schedule_interval='@daily',
    start_date=days_ago(1),
    catchup=False
) as dag:
```

- Defines the DAG (Directed Acyclic Graph)
- `schedule\_interval='@daily'`: Runs once per day
- `start\_date=days\_ago(1)`: Starts from yesterday
- `catchup=False`: Doesn't run missed schedules
- **Why**: Automates daily model retraining
- **Usage**: Change to `@hourly` for more frequent runs

```
ingest_task = BashOperator(
    task_id='ingest_data',
    bash_command='python /opt/airflow/src/ingestion.py'
)
```

- Task to load raw data
- **Why**: First step in pipeline
- **Usage**: Runs ingestion.py script

```
validate_task = BashOperator(
    task_id='validate_data',
    bash_command='python /opt/airflow/src/validate_data.py'
)
```

- Task to check data quality
- **Why**: Prevents training on corrupted data
- **Usage**: Fails pipeline if data is invalid

```
train_task = BashOperator(
    task_id='train_model',
    bash_command='python /opt/airflow/src/train.py'
)
```

- Task to train ML model
- **Why**: Core ML training step
- **Usage**: Creates new model file

```
predict_task = PythonOperator(
    task_id='test_internal_prediction',
    python_callable=run_inference_test
)
```

- Task to test model internally
- **Why**: Verifies model works before API test
- **Usage**: Calls Python function directly

```
api_task = PythonOperator(
    task_id='call_live_api_tracking',
    python_callable=call_live_api_predict
)
```

- Task to test live API
- **Why**: End-to-end validation
- **Usage**: Ensures API is serving predictions

```
ingest_task >> validate_task >> train_task >> predict_task >> api_task
```

- Defines task execution order
- `>>` means "then"
- **Why**: Creates pipeline flow
- **Usage**: Tasks run sequentially, each waits for previous to complete

### **When to Use:**

- Automatically runs daily at midnight
- Manually trigger from Airflow UI

- Modify schedule\_interval for different frequencies

### ***\*\*dags/monitoring\_dag.py\*\****

**Purpose:** Monitors model performance, detects data drift, sends alerts, and triggers retraining.

```
from evidently.report import Report
from evidently.metric_preset import DataDriftPreset
from evidently.pipeline.column_mapping import ColumnMapping
```

- Evidently AI library for drift detection
- **Why**: Detects when data distribution changes
- **Usage**: Compares current data to training data

```
def upload_to_s3(file_path, object_name):
    bucket_name = "monitoring-reports"
    s3_client = boto3.client(
        's3',
        endpoint_url='http://localstack:4566',
        aws_access_key_id='test',
        aws_secret_access_key='test',
        region_name='us-east-1'
    )
```

- Uploads reports to LocalStack S3
- **Why**: Stores monitoring reports for historical analysis
- **Usage**: Access reports from S3 bucket

```
try:
    s3_client.head_bucket(Bucket=bucket_name)
except:
    print(f"■ Creating bucket: {bucket_name}")
    s3_client.create_bucket(Bucket=bucket_name)
```

- Checks if bucket exists, creates if not
- **Why**: Auto-setup, no manual bucket creation needed
- **Usage**: First run creates bucket automatically

```
s3_client.upload_file(file_path, bucket_name, object_name)
```

- Uploads HTML report to S3
- **Why**: Persistent storage of monitoring results
- **Usage**: Download reports from LocalStack

```
def generate_monitoring_report():
    engine = create_engine("postgresql://airflow:airflow@postgres/airflow")
    query = "SELECT * FROM predictions"
    df = pd.read_sql(query, engine)
```

- Loads all predictions from database
- **Why**: Analyzes recent prediction data
- **Usage**: Compares to training data distribution

```
if len(df) < 5:
    print("■ Not enough data for monitoring!")
```

```
    return False
```

- Requires minimum 5 predictions
- **Why**: Statistical analysis needs sufficient data
- **Usage**: Wait for more predictions before monitoring

```
reference_df = df.sample(n=len(df)//2, random_state=42)
current_df = df.drop(reference_df.index)
```

- Splits data into reference and current sets
- **Why**: Compares two time periods for drift
- **Usage**: Simulates before/after comparison

```
column_mapping = ColumnMapping()
column_mapping.prediction = 'predicted_cnt'
column_mapping.numerical_features = ['temp', 'atemp', 'hum', 'windspeed']
column_mapping.categorical_features = ['season', 'mnth', 'hr', 'holiday', 'weekday', 'workingday', 'weat']
```

- Tells Evidently which columns are which type
- **Why**: Proper statistical tests for each feature type
- **Usage**: Customize based on your features

```
monitoring_report = Report(metrics=[DataDriftPreset()])
monitoring_report.run(
    reference_data=reference_df,
    current_data=current_df,
    column_mapping=column_mapping
)
```

- Generates drift detection report
- **Why**: Identifies which features have changed
- **Usage**: Review report to understand drift

```
report_path_json = "/opt/airflow/src/monitoring_report.json"
report_path_html = "/opt/airflow/src/monitoring_report.html"
monitoring_report.save_json(report_path_json)
monitoring_report.save_html(report_path_html)
```

- Saves report in two formats
- JSON: Machine-readable for automation
- HTML: Human-readable for review
- **Why**: Different formats for different uses
- **Usage**: View HTML in browser, parse JSON in code

```
report_ts = datetime.now().strftime("%Y%m%d_%H%M%S")
upload_to_s3(report_path_html, f"bike_report_{report_ts}.html")
```

- Adds timestamp to filename
- **Why**: Keeps history of all reports
- **Usage**: Compare reports over time

```
def check_for_issues(ti):
    report_path_json = "/opt/airflow/logs/monitoring_report.json"
    if not os.path.exists(report_path_json):
```

```
        return False
```

- Checks if drift was detected
- `ti`: Task Instance (Airflow context)
- \*\*Why\*\*: Decides whether to trigger alerts
- \*\*Usage\*\*: Returns True if action needed

```
with open(report_path_json, 'r') as f:  
    data = json.load(f)  
  
drift_detected = False  
for metric in data.get('metrics', []):  
    res = metric.get('result', {})  
    if 'dataset_drift' in res:  
        drift_detected = res['dataset_drift']
```

- Parses JSON report
- Extracts drift status
- \*\*Why\*\*: Automated decision making
- \*\*Usage\*\*: True if significant drift found

```
ti.xcom_push(key='drift_status', value="DETECTED █" if drift_detected else "Normal █")
```

- Stores drift status in XCom (Airflow's inter-task communication)
- \*\*Why\*\*: Passes data to downstream tasks
- \*\*Usage\*\*: Slack alert task reads this value

```
if drift_detected:  
    print("█ Drift found! Triggering alerts and retraining.")  
    return True  
else:  
    print("█ Everything is normal. Skipping retraining.")  
    return False
```

- Returns True only if drift detected
- \*\*Why\*\*: ShortCircuitOperator stops pipeline if False
- \*\*Usage\*\*: Saves resources by skipping unnecessary retraining

```
def send_slack_manual(ti):  
    webhook_url = "https://hooks.slack.com/services/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"  
    drift = ti.xcom_pull(task_ids='check_for_issues', key='drift_status')
```

- Sends alert to Slack channel
- Retrieves drift status from XCom
- \*\*Why\*\*: Notifies team of issues
- \*\*Usage\*\*: Replace webhook\_url with your Slack webhook

```
payload = {  
    "text": (  
        "█ *MLOps Alert: Action Required!* \n\n"  
        f"• *Status:* {drift}\n"  
        "• *Reason:* ████ ████ ████ (Drift), ████ ████ *Retrain* ████  
    )  
}  
requests.post(webhook_url, json=payload, timeout=10)
```

- Formats and sends Slack message
- **Why**: Human-readable alert with context
- **Usage**: Customize message format

```
with DAG(
    dag_id='model_monitoring_dag',
    start_date=datetime(2026, 1, 1),
    schedule_interval='@daily',
    catchup=False
) as dag:
```

- Monitoring DAG runs daily
- **Why**: Regular checks for model degradation
- **Usage**: Adjust schedule based on data volume

```
monitor_task = PythonOperator(task_id='generate_report', python_callable=generate_monitoring_report)
check_issues_task = ShortCircuitOperator(task_id='check_for_issues', python_callable=check_for_issues)
send_slack_task = PythonOperator(task_id='send_slack_alert', python_callable=send_slack_manual)
retrain_task = PythonOperator(task_id='retrain_model', python_callable=run_pipeline)
```

- Defines monitoring tasks
- `ShortCircuitOperator`: Stops pipeline if returns False
- **Why**: Conditional execution based on drift
- **Usage**: Only retrains if necessary

```
monitor_task >> check_issues_task >> [send_slack_task, retrain_task]
```

- Pipeline flow with branching
- `[send\_slack\_task, retrain\_task]`: Both run in parallel
- **Why**: Alert and retrain happen simultaneously
- **Usage**: Efficient use of resources

### **When to Use:**

- Runs automatically daily
- Manually trigger after significant data changes
- Review HTML reports to understand drift patterns

### **\*\*dags/test\_dags.py\*\***

**Purpose:** Simple test DAG to verify Airflow is working.

```
def hello():
    print("Hello Rama!")
```

- Minimal test function
- **Why**: Simplest possible task
- **Usage**: Debugging Airflow setup

```
with DAG('simple_test_dag', start_date=datetime(2026, 1, 1), schedule='@daily', catchup=False) as dag:
    t1 = PythonOperator(task_id='say_hello', python_callable=hello)
```

- Creates basic DAG with one task
- **Why**: Verifies Airflow can execute Python code
- **Usage**: First DAG to test after setup

#### **When to Use:**

- After fresh Airflow installation
- Troubleshooting DAG loading issues
- Learning Airflow basics

### **3. SOURCE CODE FILES (src/)**

#### **\*\*src/api.py\*\***

**Purpose:** FastAPI server that serves predictions and logs them to PostgreSQL.

```
from prometheus_fastapi_instrumentator import Instrumentator
```

- Prometheus integration for FastAPI
- **Why**: Automatic metrics collection (request count, latency, etc.)
- **Usage**: Metrics available at `/metrics` endpoint

```
app = FastAPI(title="BikesBooking Prediction API")
Instrumentator().instrument(app).expose(app)
```

- Creates FastAPI application
- Instruments it with Prometheus
- **Why**: Enables monitoring without manual code
- **Usage**: Prometheus scrapes `/metrics` automatically

```
DATABASE_URL = "postgresql+psycopg2://airflow:airflow@localhost:5432/airflow"
engine = create_engine(DATABASE_URL)
SessionLocal = sessionmaker(autocommit=False, autoflush=False, bind=engine)
Base = declarative_base()
```

- Database connection setup
- `localhost:5432`: API runs on Windows, connects to Docker PostgreSQL
- **Why**: Stores all predictions for monitoring
- **Usage**: Every prediction is logged

```
class PredictionLog(Base):
    __tablename__ = "predictions"
    id = Column(Integer, primary_key=True, index=True)
    timestamp = Column(DateTime, default=datetime.utcnow)
    season = Column(Integer); mnth = Column(Integer); hr = Column(Integer)
    holiday = Column(Integer); weekday = Column(Integer); workingday = Column(Integer)
    weathersit = Column(Integer); temp = Column(Float); atemp = Column(Float)
    hum = Column(Float); windspeed = Column(Float)
    predicted_cnt = Column(Float)
```

- SQLAlchemy ORM model for predictions table
- **Why**: Type-safe database operations
- **Usage**: Automatically creates table if not exists

```
Base.metadata.create_all(bind=engine)
```

- Creates database tables
- **Why**: Auto-setup on first run
- **Usage**: No manual SQL needed

```
@app.get("/")
def home():
    return {
        "message": "API Live",
        "db": "Postgres Connected",
        "monitoring": "Prometheus Enabled at /metrics"
    }
```

- Health check endpoint
- **Why**: Verify API is running
- **Usage**: Visit `http://localhost:9999/` in browser

```
@app.get("/predict")
async def predict(
    season: int = 1, mnth: int = 1, hr: int = 12,
    holiday: int = 0, weekday: int = 0, workingday: int = 1,
    weathersit: int = 1, temp: float = 0.5, atemp: float = 0.5,
    hum: float = 0.5, windspeed: float = 0.0
):
```

- Prediction endpoint with query parameters
- Default values provided
- **Why**: Easy to test in browser
- **Usage**: `http://localhost:9999/predict?temp=0.7&hr=18`

```
db = SessionLocal()
try:
    data_dict = {
        "season": season, "mnth": mnth, "hr": hr,
        "holiday": holiday, "weekday": weekday, "workingday": workingday,
        "weathersit": weathersit, "temp": temp, "atemp": atemp,
        "hum": hum, "windspeed": windspeed
    }
```

- Converts parameters to dictionary
- **Why**: Format expected by prediction function
- **Usage**: Matches model training features

```
prediction = make_prediction(data_dict)
predicted_val = max(0, int(round(float(prediction[0]))))
```

- Makes prediction using loaded model
- Ensures non-negative integer result
- **Why**: Can't have negative bikes
- **Usage**: Returns realistic bike count

```

new_log = PredictionLog(**data_dict, predicted_cnt=float(predicted_val))
db.add(new_log)
db.commit()
print(f"■ Saved to Postgres: {predicted_val}")

```

- Saves prediction to database
- \*\*Why\*\*: Enables monitoring and drift detection
- \*\*Usage\*\*: All predictions tracked automatically

```

        return {"predicted_bikes": predicted_val}
    except Exception as e:
        print(f"■ API Error: {str(e)}")
        return {"error": str(e)}
    finally:
        db.close()

```

- Returns prediction or error
- Always closes database connection
- \*\*Why\*\*: Prevents connection leaks
- \*\*Usage\*\*: Graceful error handling

```

if __name__ == "__main__":
    uvicorn.run(app, host="0.0.0.0", port=9999)

```

- Starts API server
- `0.0.0.0`: Accepts connections from any IP
- \*\*Why\*\*: Allows Docker containers to connect
- \*\*Usage\*\*: Run `python src/api.py` on Windows

### **When to Use:**

- Run on Windows host: `python src/api.py`
- Access from browser: `http://localhost:9999/predict?temp=0.8`
- Monitor metrics: `http://localhost:9999/metrics`

### **\*\*src/app.py\*\***

**Purpose:** Streamlit web interface for predictions and monitoring.

```

st.set_page_config(page_title="Bike Sharing MLOps Portal", layout="wide")

```

- Configures Streamlit page
- `layout="wide"`: Uses full browser width
- \*\*Why\*\*: Better visualization of charts
- \*\*Usage\*\*: Must be first Streamlit command

```

st.title("■ Bike Sharing Prediction & Monitoring")
st.markdown("---")

```

- Page title and separator

- **Why**: Professional appearance

- **Usage**: Customize branding

```
col1, col2 = st.columns([1, 2])
```

- Creates two-column layout
- Ratio 1:2 (sidebar smaller than main)
- **Why**: Organized interface
- **Usage**: Inputs on left, results on right

```
with col1:
    st.header("■■■ Input Features")
    season = st.selectbox("Season", [1, 2, 3, 4], help="1:Spring, 2:Summer, 3:Fall, 4:Winter")
    mnth = st.slider("Month", 1, 12, 6)
    hr = st.slider("Hour of Day", 0, 23, 12)
    temp = st.slider("Temperature (Normalized)", 0.0, 1.0, 0.5)
    hum = st.slider("Humidity (Normalized)", 0.0, 1.0, 0.5)
```

- Input widgets for features
- **Why**: User-friendly data entry
- **Usage**: Adjust sliders to see different predictions

```
predict_btn = st.button("■ Predict Bike Demand", use_container_width=True)
```

- Prediction trigger button
- `use\_container\_width`: Full width button
- **Why**: Clear call-to-action
- **Usage**: Click to get prediction

```
with col2:
    st.header("■ Prediction Results")
    if predict_btn:
        url = "http://host.docker.internal:9999/predict"
```

- Results display area
- Only executes when button clicked
- **Why**: Prevents unnecessary API calls
- **Usage**: Shows prediction after button click

```
params = {
    "season": season, "mnth": mnth, "hr": hr,
    "holiday": 0, "weekday": 3, "workingday": 1,
    "weathersit": 1, "temp": temp, "atemp": atemp,
    "hum": hum, "windspeed": 0.1
}
```

- Builds API request parameters
- Some values hardcoded (holiday, weekday)
- **Why**: Simplifies UI, focuses on key features
- **Usage**: Extend UI to include all features

```
try:
    with st.spinner('Calculating demand...'):
        response = requests.get(url, params=params, timeout=5)
```

- Shows loading spinner during API call
- 5 second timeout
- **Why**: User feedback during wait
- **Usage**: Prevents UI freeze

```
if response.status_code == 200:
    res = response.json()
    prediction = res.get('predicted_bikes', 0)
    st.metric(label="Predicted Bikes Needed", value=int(prediction))
    st.balloons()
```

- Displays prediction result
- Shows balloons animation on success
- **Why**: Positive user feedback
- **Usage**: Clear result presentation

```
else:
    st.error(f"■ API Error: {response.text}")
except Exception as e:
    st.error(f"■ Connection Failed: Make sure api.py is running on port 9999!")
```

- Error handling with user-friendly messages
- **Why**: Helps troubleshoot issues
- **Usage**: Guides user to fix problems

```
st.header("■ Model Monitoring")
st.info("View the latest Data Drift analysis from Evidently AI.")

if st.button("■ View Data Drift Report"):
    report_path = "monitoring_report.html"
```

- Monitoring section
- **Why**: Integrated monitoring in same UI
- **Usage**: Click to view drift report

```
if os.path.exists(report_path):
    try:
        with open(report_path, 'r', encoding='utf-8') as f:
            html_data = f.read()
            components.html(html_data, height=1000, scrolling=True)
```

- Embeds HTML report in Streamlit
- **Why**: No need to open separate file
- **Usage**: Scroll through report in app

```
else:
    st.warning(f"■■ Report file not found at: {report_path}")
    st.info("Please run the Airflow DAG first to generate the report.")
```

- Helpful message if report missing
- **Why**: Guides user to generate report
- **Usage**: Run monitoring DAG first

```

st.sidebar.markdown(f"""
---  

**System Status:**  

- **Frontend:** Running (Docker)  

- **Shared Path:** `/opt/airflow/src/  

- **Backend API:** http://localhost:9999  

- **Database:** PostgreSQL  

- **Monitoring:** Evidently AI
""")
```

- System information in sidebar
- **Why**: Quick reference for debugging
- **Usage**: Check configuration at a glance

### **When to Use:**

- Access at `http://localhost:8501`
- Make predictions interactively
- View monitoring reports

### ***src/config.yaml***

**Purpose:** Central configuration file for all paths and settings.

```

data:  

  url: "https://raw.githubusercontent.com/arashgmn/bike-sharing/master/hour.csv"
```

- Source URL for raw data
- **Why**: Single source of truth for data location
- **Usage**: Change URL to use different dataset

```

raw_path: "/opt/airflow/data/bike_sharing_raw.csv"
processed_dir: "/opt/airflow/data/processed"
test_size: 0.2
```

- File paths and train/test split ratio
- **Why**: Consistent paths across all scripts
- **Usage**: Modify paths for different environments

```

model:  

  save_path: "/opt/airflow/models/bike_model.pkl"
```

- Model file location
- **Why**: All scripts save/load from same location
- **Usage**: Change path for model versioning

```

aws:  

  endpoint_url: "http://localstack:4566"  

  bucket_name: "bike-sharing"  

  region_name: "us-east-1"
```

- LocalStack S3 configuration
- **Why**: Simulates AWS without costs

- \*\*Usage\*\*: Replace with real AWS credentials for production

#### **When to Use:**

- Modify before running pipeline
- Different configs for dev/staging/prod
- Version control configuration changes

#### ***\*\*src/ingestion.py\*\****

**Purpose:** Loads raw data for the pipeline.

```
def load_data():
    data_file = "/opt/airflow/data/bike_sharing_raw.csv"
    print(f"■ Checking for local data file at: {data_file}")
```

- Checks if data file exists
- \*\*Why\*\*: Verifies data availability before processing
- \*\*Usage\*\*: First step in pipeline

```
if os.path.exists(data_file):
    print(f"■ Found it! Data is already present. Skipping download.")
    size = os.path.getsize(data_file)
    print(f"■ File size: {size} bytes")
    sys.exit(0)
```

- Confirms file exists and shows size
- `sys.exit(0)`: Success exit code
- \*\*Why\*\*: Airflow marks task as successful
- \*\*Usage\*\*: No download needed if file present

```
else:
    print(f"■ Error: File NOT found at {data_file}!")
    sys.exit(1)
```

- Fails if file missing
- `sys.exit(1)`: Error exit code
- \*\*Why\*\*: Stops pipeline early if data missing
- \*\*Usage\*\*: Prevents training on missing data

#### **When to Use:**

- First task in Airflow DAG
- Manually: `python src/ingestion.py`
- Extend to download from URL if missing

#### ***\*\*src/preprocessing.py\*\****

**Purpose:** Cleans and splits data for training.

```
def preprocess():
    with open("config.yaml", "r") as f:
        config = yaml.safe_load(f)
```

- Loads configuration
- **Why**: Uses centralized settings
- **Usage**: Consistent with other scripts

```
raw_data_path = config['data']['raw_path']
df = pd.read_csv(raw_data_path)
```

- Loads raw data
- **Why**: Reads from configured path
- **Usage**: Flexible data location

```
df = df.drop(columns= ['instant', 'dteday', 'casual', 'registered'])
```

- Removes unnecessary columns
- `instant`: Row ID (not a feature)
- `dteday`: Date string (already encoded in other features)
- `casual`, `registered`: Components of target variable
- **Why**: Prevents data leakage and reduces noise
- **Usage**: Keep only predictive features

```
X = df.drop(columns=['cnt'])
y = df['cnt']
```

- Separates features (X) from target (y)
- `cnt`: Total bike rentals (what we predict)
- **Why**: Standard ML practice
- **Usage**: X for input, y for output

```
test_size = config['data']['test_size']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size, random_state=42)
```

- Splits data 80/20 (train/test)
- `random\_state=42`: Reproducible splits
- **Why**: Evaluate model on unseen data
- **Usage**: Test set measures generalization

```
processed_dir = config['data']['processed_dir']
os.makedirs(processed_dir, exist_ok=True)
```

- Creates output directory
- `exist\_ok=True`: No error if already exists
- **Why**: Ensures directory exists before saving
- **Usage**: Automatic directory creation

```
X_train.to_csv(f"{processed_dir}/X_train.csv", index=False)
X_test.to_csv(f"{processed_dir}/X_test.csv", index=False)
```

```
y_train.to_csv(f"{processed_dir}/y_train.csv" , index=False)
y_test.to_csv(f"{processed_dir}/y_test.csv" , index=False)
```

- Saves processed data
- `index=False`: Don't save row numbers
- \*\*Why\*\*: Separate files for each dataset
- \*\*Usage\*\*: Training script loads these files

### When to Use:

- After data ingestion
- Before model training
- Manually: `python src/preprocessing.py`

### \*\*src/train.py\*\*

**Purpose:** Trains Random Forest model and registers it in MLflow.

```
def train_model():
    if os.path.exists("/opt/airflow"):
        TRACKING_URI = "http://mlflow_server:5000"
        BASE_PATH = "/opt/airflow"
    else:
        TRACKING_URI = "http://localhost:5000"
        BASE_PATH = ".."
```

- Detects environment (Docker vs local)
- \*\*Why\*\*: Works in both Airflow and local development
- \*\*Usage\*\*: Automatic environment detection

```
mlflow.set_tracking_uri(TRACKING_URI)
mlflow.set_experiment("bike-sharing-experiment")
```

- Connects to MLflow server
- Creates/uses experiment
- \*\*Why\*\*: Organizes all training runs
- \*\*Usage\*\*: View experiments at `http://localhost:5000`

```
X_train_path = os.path.join(BASE_PATH, "data/processed/X_train.csv")
y_train_path = os.path.join(BASE_PATH, "data/processed/y_train.csv")
save_path = os.path.join(BASE_PATH, "models/bike_model.pkl")
```

- Defines file paths
- \*\*Why\*\*: Environment-agnostic paths
- \*\*Usage\*\*: Works in Docker and Windows

```
try:
    X_train_raw = pd.read_csv(X_train_path)
    y_train = pd.read_csv(y_train_path)
    features = ['season', 'mnth', 'hr', 'holiday', 'weekday', 'workingday', 'weathersit', 'temp', 'atemp',
    X_train = X_train_raw[[col for col in features if col in X_train_raw.columns]]
```

- Loads training data
- Selects specific features
- **Why**: Ensures consistent feature set
- **Usage**: Prevents feature mismatch errors

```
with mlflow.start_run(run_name="Airflow_Training_Run") as run:
    model = RandomForestRegressor(n_estimators=random.randint(100, 200), random_state=42)
    model.fit(X_train, y_train.values.ravel())
```

- Starts MLflow run
- Trains Random Forest model
- `n\_estimators`: Random between 100-200
- **Why**: Tracks each training run
- **Usage**: Compare different runs in MLflow UI

```
rmse = np.sqrt(mean_squared_error(y_train, model.predict(X_train)))
mlflow.log_metric("rmse", rmse)
mlflow.log_param("n_estimators", model.n_estimators)
```

- Calculates training error
- Logs metrics and parameters to MLflow
- **Why**: Track model performance over time
- **Usage**: Compare RMSE across runs

```
joblib.dump(model, "bike_model.pkl")
mlflow.log_artifact("bike_model.pkl", artifact_path="model")
```

- Saves model as artifact
- **Why**: Stores model with experiment
- **Usage**: Download model from MLflow UI

```
model_name = "Bike_Sharing_Model"
run_id = run.info.run_id
model_uri = f"runs:{run_id}/model"

print(f"■ Registering model from {model_uri}...")
```

- Prepares model registration
- **Why**: Model registry for versioning
- **Usage**: Track production models

```
client = MlflowClient()
try:
    client.create_registered_model(model_name)
except:
    pass
```

- Creates model in registry
- Ignores error if already exists
- **Why**: First-time setup
- **Usage**: One model name, multiple versions

```

res = client.create_model_version(name=model_name, source=model_uri, run_id=run_id)
current_version = res.version

client.set_registered_model_alias(model_name, "champion", str(current_version))

```

- Creates new model version
- Sets "champion" alias to latest version
- **Why**: Easy reference to production model
- **Usage**: Load model by alias instead of version number

```

os.makedirs(os.path.dirname(save_path), exist_ok=True)
joblib.dump(model, save_path)

```

- Saves model locally
- **Why**: API loads from local file
- **Usage**: Fast predictions without MLflow

```

print(f"■ SUCCESS! RMSE: {rmse:.4f} | Version: {current_version}")

```

- Success message with metrics
- **Why**: Confirms training completed
- **Usage**: Check Airflow logs for RMSE

```

except Exception as e:
    print(f"■ ERROR: {str(e)}")
    sys.exit(1)

```

- Error handling
- **Why**: Fails Airflow task on error
- **Usage**: Debug from error message

### **When to Use:**

- Airflow training task
- Manually: `python src/train.py`
- After data preprocessing

### **src/predict.py**

**Purpose:** Loads model and makes predictions.

```

def make_prediction(input_data):
    BASE_DIR = os.path.dirname(os.path.dirname(os.path.abspath(__file__)))
    model_path = os.path.join(BASE_DIR, "models", "bike_model.pkl")

```

- Finds model file relative to script location
- **Why**: Works from any directory
- **Usage**: Flexible path resolution

```

if not os.path.exists(model_path):
    print(f"■ Model not found at: {model_path}")

```

```
    return [0.0]
```

- Checks if model exists
- Returns 0 if missing
- **Why**: Graceful degradation
- **Usage**: Prevents crashes

```
model = joblib.load(model_path)
```

- Loads trained model
- **Why**: Fast loading from pickle file
- **Usage**: Model ready for predictions

```
if hasattr(model, 'monotonic_cst'):
    delattr(model, 'monotonic_cst')
```

- Removes scikit-learn version-specific attribute
- **Why**: Compatibility across versions
- **Usage**: Prevents version mismatch errors

```
try:
    input_df = pd.DataFrame([input_data]) if isinstance(input_data, dict) else pd.DataFrame(input_data)
```

- Converts input to DataFrame
- Handles both dict and DataFrame inputs
- **Why**: Flexible input format
- **Usage**: Works with API and direct calls

```
cols = ['season', 'mnth', 'hr', 'holiday', 'weekday', 'workingday', 'weathersit', 'temp', 'atemp',
final_df = input_df.reindex(columns=cols, fill_value=0)
final_df = final_df.apply(pd.to_numeric)
```

- Ensures correct column order
- Fills missing columns with 0
- Converts to numeric types
- **Why**: Matches training data format
- **Usage**: Prevents feature order errors

```
prediction = model.predict(final_df)
return [float(prediction[0])]
```

- Makes prediction
- Returns as list of floats
- **Why**: Consistent return format
- **Usage**: Easy to parse in API

```
except Exception as e:
    print(f"■ Prediction Logic Error: {e}")
    return [0.0]
```

- Error handling
- Returns 0 on error
- **Why**: API doesn't crash
- **Usage**: Check logs for errors

### **When to Use:**

- Called by API for each prediction
- Called by Airflow for testing
- Import in other scripts: `from predict import make\_prediction`

### **\*\*src/evaluate.py\*\***

**Purpose:** Evaluates model performance on test set.

```
def evaluate_model():
    with open("config.yaml", "r") as f:
        config = yaml.safe_load(f)

    processed_dir = config['data']['processed_dir']
    X_test = pd.read_csv(f"{processed_dir}/X_test.csv")
    y_test = pd.read_csv(f"{processed_dir}/y_test.csv")
```

- Loads configuration
- **Why**: Uses centralized paths
- **Usage**: Consistent with other scripts

```
model_path = config['model']['save_path']
model = joblib.load(model_path)
```

- Loads test data
- **Why**: Evaluates on unseen data
- **Usage**: Measures generalization

```
predictions = model.predict(X_test)
```

- Makes predictions on test set
- **Why**: Compare to actual values
- **Usage**: Calculates metrics

```
mse = mean_squared_error(y_test, predictions)
r2 = r2_score(y_test, predictions)
```

- Calculates evaluation metrics
- MSE: Mean Squared Error (lower is better)

- **R<sup>2</sup>**: Coefficient of determination (higher is better, max 1.0)
- **Why**: Quantifies model performance
- **Usage**: Compare models

```
print("\n---- Model Report-----")
print(f"Mean Squared Error: {mse:.2f}")
print(f"R2 Score: {r2:.2f}")
```

- Prints metrics
- **Why**: Easy to read results
- **Usage**: Check console output

### **When to Use:**

- After training
- Manually: `python src/evaluate.py`
- Add to Airflow DAG for automated evaluation

### **\*\*src/validate\_data.py\*\***

**Purpose:** Checks data quality before training.

```
def validate_data():
    path = "/opt/airflow/data/bike_sharing_raw.csv"

    if not os.path.exists(path):
        print(f"■ Error: File {path} not found!")
        sys.exit(1)
```

- Checks if data file exists
- **Why**: Prevents training on missing data
- **Usage**: Early failure detection

```
df = pd.read_csv(path)
print(f"■ Data loaded. Columns: {list(df.columns)}")
```

- Loads data and shows columns
- **Why**: Verifies file is readable
- **Usage**: Debug column name issues

```
errors = []

if 'cnt' not in df.columns:
    errors.append("Column 'cnt' missing")
elif df['cnt'].isnull().any():
    errors.append("Null values found in 'cnt'")
```

- Checks for target column
- Checks for missing values
- **Why**: Target is required for training
- **Usage**: Ensures data quality

```
if 'temp' not in df.columns:
```

```
errors.append("Column 'temp' missing")
```

- Checks for key feature
- \*\*Why\*\*: Model needs temperature data
- \*\*Usage\*\*: Add more checks for other features

```
if not errors:  
    print("■ Data Validation SUCCESSFUL (via Pandas)!")  
    sys.exit(0)  
else:  
    print(f"■ Data Validation FAILED: {errors}")  
    sys.exit(1)
```

- Reports validation result
- \*\*Why\*\*: Stops pipeline if data is bad
- \*\*Usage\*\*: Airflow marks task as failed

### When to Use:

- After data ingestion
- Before preprocessing
- Extend with more validation rules

### \*\*src/run\_pipeline.py\*\*

**Purpose:** Standalone training script for retraining.

```
def run_pipeline():  
    print("\n■ Real Data Training Started...")  
    csv_path = r'C:\MLOPS\data\bike_sharing_raw.csv'
```

- Trains model from scratch
- \*\*Why\*\*: Used by monitoring DAG for retraining
- \*\*Usage\*\*: Called when drift detected

```
if not os.path.exists(csv_path):  
    print(f"■ File Not Found at: {csv_path}")  
    return
```

- Checks for data file
- \*\*Why\*\*: Graceful failure
- \*\*Usage\*\*: Returns without error

```
df = pd.read_csv(csv_path)  
  
expected_cols = [  
    'season', 'yr', 'mnth', 'hr', 'holiday', 'weekday',  
    'workingday', 'weathersit', 'temp', 'atemp', 'hum', 'windspeed'  
]  
  
X = df[expected_cols]  
y = df['cnt']
```

- Loads data and selects features

- **Why**: Self-contained training
- **Usage**: No dependency on preprocessing script

```
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

- Splits data
- **Why**: Evaluates new model
- **Usage**: Standard train/test split

```
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(x_train, y_train)
```

- Trains model
- **Why**: Creates new model version
- **Usage**: Replaces old model

```
accuracy = model.score(x_test, y_test)
print(f"■ New Model Accuracy: {accuracy:.4f}")
```

- Calculates R<sup>2</sup> score
- **Why**: Verifies new model quality
- **Usage**: Compare to old model

```
joblib.dump(model, 'bike_model.pkl')
print("■ Model Saved as 'bike_model.pkl'!")
```

- Saves new model
- **Why**: Replaces production model
- **Usage**: API uses new model immediately

### **When to Use:**

- Called by monitoring DAG
- Manually: `python src/run\_pipeline.py`
- When drift detected

### ***src/drift\_detection.py***

**Purpose:** Detects data drift using Evidently AI.

```
def check_model_drift(reference_path, current_data_df):
    reference_df = pd.read_csv(reference_path)
```

- Loads reference data (training data)
- **Why**: Baseline for comparison
- **Usage**: Pass path to training data

```
drift_report = Report(metrics=[DataDriftPreset()])
```

- Creates Evidently report
- **Why**: Pre-configured drift metrics
- **Usage**: Comprehensive drift analysis

```
drift_report.run(reference_data=reference_df, current_data=current_data_df)
```

- Compares datasets
- **Why**: Statistical tests for each feature
- **Usage**: Detects distribution changes

```
report_dict = drift_report.as_dict()
drift_detected = report_dict['metrics'][0]['result']['dataset_drift']
```

- Extracts drift status
- **Why**: Boolean decision for automation
- **Usage**: True if significant drift

```
return drift_detected, report_dict
```

- Returns drift status and full report
- **Why**: Both automated and manual analysis
- **Usage**: Use boolean for alerts, dict for details

### **When to Use:**

- Import in monitoring scripts
- Customize drift thresholds
- Add to monitoring DAG

### ***\*\*src/monitor.py\*\****

**Purpose:** Monitors model performance and triggers retraining.

```
def monitor_and_retrain():
    print("\n■ Step 1: Loading model and data for monitoring...")

    model_path = 'bike_model.pkl'
    csv_path = r'C:\MLOPS\data\bike_sharing_raw.csv'
```

- Loads model and data
- **Why**: Checks current model performance
- **Usage**: Run periodically

```
if not os.path.exists(model_path):
    print("■ Model not found! Starting training...")
    run_pipeline()
    return
```

- Trains model if missing
- **Why**: Bootstrap initial model

- **\*\*Usage\*\*:** First-time setup

```
model = joblib.load(model_path)
df = pd.read_csv(csv_path)
```

- Loads model and data
- **\*\*Why\*\*:** Evaluates current performance
- **\*\*Usage\*\*:** Uses latest data

```
expected_cols = [
    'season', 'yr', 'mnth', 'hr', 'holiday', 'weekday',
    'workingday', 'weathersit', 'temp', 'atemp', 'hum', 'windspeed'
]

X = df[expected_cols]
y_true = df['cnt']
```

- Prepares data
- **\*\*Why\*\*:** Matches training format
- **\*\*Usage\*\*:** Consistent feature set

```
print("■ Calculating current accuracy...")
y_pred = model.predict(X)
current_score = r2_score(y_true, y_pred)
print(f"■ Current Accuracy: {current_score:.4f}")
```

- Calculates current R<sup>2</sup> score
- **\*\*Why\*\*:** Measures model degradation
- **\*\*Usage\*\*:** Compare to initial performance

```
THRESHOLD = 0.92

if current_score < THRESHOLD:
    print(f"■■■ Accuracy is below threshold {THRESHOLD}. Triggering Retraining...")
    run_pipeline()
    print("■ Success: Model has been retrained and updated!")
else:
    print("■ Accuracy is good. No action needed.")
```

- Compares to threshold
- Retrains if below threshold
- **\*\*Why\*\*:** Automated model maintenance
- **\*\*Usage\*\*:** Adjust threshold based on requirements

### **When to Use:**

- Run manually: `python src/monitor.py`
- Schedule with cron or Task Scheduler
- Integrate with monitoring DAG

### **\*\*src/requirements.txt\*\***

**Purpose:** Lists all Python dependencies.

```
fastapi
```

uvicorn

- FastAPI framework and ASGI server
- **Why**: API serving
- **Usage**: `pip install fastapi uvicorn`

mlflow

- Experiment tracking and model registry
- **Why**: ML lifecycle management
- **Usage**: Track experiments

pandas==1.5.3  
scikit-learn==1.3.2

- Data manipulation and ML library
- Specific versions for compatibility
- **Why**: Core ML functionality
- **Usage**: Data processing and modeling

joblib  
pyyaml

- Model serialization and config parsing
- **Why**: Save/load models and configs
- **Usage**: Persistence

great\_expectations==0.16.15  
evidently==0.4.15

- Data validation and drift detection
- **Why**: Data quality and monitoring
- **Usage**: Automated checks

sqlalchemy  
psycopg2-binary

- Database ORM and PostgreSQL driver
- **Why**: Database operations
- **Usage**: Store predictions

pydantic==1.10.13

- Data validation
- Specific version for compatibility
- **Why**: FastAPI dependency
- **Usage**: Request validation

prometheus-fastapi-instrumentator

- Prometheus metrics for FastAPI
- \*\*Why\*\*: API monitoring
- \*\*Usage\*\*: Automatic metrics

```
streamlit
requests
```

- Web UI framework and HTTP client
- \*\*Why\*\*: Frontend and API calls
- \*\*Usage\*\*: User interface

**When to Use:**

- Install all: `pip install -r src/requirements.txt`
- Update versions as needed
- Freeze versions: `pip freeze > requirements.txt`

## 4. CONFIGURATION FILES

### *\*\*prometheus.yml\*\**

**Purpose:** Configures Prometheus to scrape metrics from the API.

```
global:
  scrape_interval: 5s
```

- How often to collect metrics
- \*\*Why\*\*: 5 seconds for near real-time monitoring
- \*\*Usage\*\*: Increase for less frequent updates

```
scrape_configs:
  - job_name: 'prediction_api'
```

- Defines a monitoring job
- \*\*Why\*\*: Groups related metrics
- \*\*Usage\*\*: Add more jobs for other services

```
static_configs:
  - targets: ['host.docker.internal:9999']
```

- Where to collect metrics from
- \*\*Why\*\*: API runs on Windows host
- \*\*Usage\*\*: Prometheus scrapes `/metrics` endpoint

**When to Use:**

- Automatically loaded by Prometheus container

- Modify to add more services
- View metrics at `http://localhost:9090`

### ***\*\*.gitignore\*\****

**Purpose:** Tells Git which files to ignore.

```
venv/
__pycache__/
```

- Python virtual environment and cache
- **Why**: Environment-specific, not needed in repo
- **Usage**: Automatic

```
mlops/
```

- MLflow experiment data
- **Why**: Large files, regenerated locally
- **Usage**: Each developer has own experiments

```
data/raw/*.csv
data/processed/*.csv
```

- Data files
- **Why**: Large files, use DVC instead
- **Usage**: Track data with DVC

```
models/*.pkl
```

- Model files
- **Why**: Large binary files
- **Usage**: Store in MLflow or DVC

```
.DS_Store
.env
```

- Mac system files and environment variables
- **Why**: System-specific and secrets
- **Usage**: Never commit secrets

### **When to Use:**

- Automatically used by Git
- Add patterns for new file types
- Keep secrets out of version control

### ***\*\*.dvcignore\*\****

**Purpose:** Tells DVC which files to ignore.

```
# Add patterns of files dvc should ignore, which could improve
# the performance. Learn more at
# https://dvc.org/doc/user-guide/dvcignore
```

- Currently empty (uses defaults)
- **Why:** DVC ignores .git, .dvc, etc. by default
- **Usage:** Add patterns if needed

**When to Use:**

- Exclude temporary files from DVC tracking
- Improve DVC performance
- Similar to .gitignore but for DVC

### ***\*\*.dvc/config\*\****

**Purpose:** DVC configuration for data versioning.

```
[core]
remote = local_storage
```

- Default remote storage
- **Why:** Where to push/pull data
- **Usage:** DVC commands use this remote

```
['remote "local_storage"']
url = C:\dvc_storage
```

- Local directory for data storage
- **Why:** Simulates remote storage locally
- **Usage:** Change to S3/GCS for production

**When to Use:**

- Automatically used by DVC commands
- Configure remote storage
- `dvc push` uploads to this location

## **5. CI/CD FILES**

### ***\*\*.github/workflows/main.yml\*\****

**Purpose:** GitHub Actions workflow for automated testing and deployment.

```
name: Bike-Sharing MLOps CI/CD
```

- Workflow name
- **Why**: Identifies workflow in GitHub UI
- **Usage**: Shows in Actions tab

```
on:
  push:
    branches: [ main ]
  pull_request:
    branches: [ main ]
```

- Triggers workflow on push or PR to main
- **Why**: Automated testing on code changes
- **Usage**: Runs automatically

```
jobs:
  test-and-lint:
    runs-on: ubuntu-latest
```

- First job: code quality checks
- **Why**: Catch errors before deployment
- **Usage**: Runs on GitHub servers

```
steps:
  - name: Checkout Code
    uses: actions/checkout@v3
```

- Downloads repository code
- **Why**: Needed to run tests
- **Usage**: Standard first step

```
- name: Set up Python
  uses: actions/setup-python@v4
  with:
    python-version: '3.9'
```

- Installs Python 3.9
- **Why**: Matches development environment
- **Usage**: Consistent Python version

```
- name: Install Linting Tools
  run: |
    pip install flake8
```

- Installs code linter
- **Why**: Checks code style
- **Usage**: Enforces standards

```
- name: Lint with flake8
  run: |
    flake8 . --count --select=E9,F63,F7,F82 --show-source --statistics
    flake8 . --count --exit-zero --max-complexity=10 --max-line-length=127 --statistics
```

- Runs linting checks
- First line: Syntax errors (fails build)

- Second line: Style warnings (doesn't fail)
- **Why**: Maintains code quality
- **Usage**: Fix errors before merging

```
build-check:
  needs: test-and-lint
  runs-on: ubuntu-latest
```

- Second job: build verification
- `needs`: Runs only if linting passes
- **Why**: Ensures Docker image builds
- **Usage**: Catches build errors early

```
steps:
  - name: Checkout Code
    uses: actions/checkout@v3

  - name: Build Docker Image (Simulated CD)
    run: |
      echo "■ Starting Docker Build to verify Cloud Readiness..."
      docker build -t bike-sharing-mlops:${{ github.sha }} .
      echo "■ Build Successful! This image is ready for AWS/Azure deployment."
```

- Builds Docker image
- Tags with commit SHA
- **Why**: Verifies deployment readiness
- **Usage**: Catches Docker errors

```
- name: Success Message
  run: echo "Final CI/CD Check Passed. The pipeline is robust and production-ready!"
```

- Success confirmation
- **Why**: Clear feedback
- **Usage**: Shows in GitHub Actions log

### **When to Use:**

- Automatically runs on every push
- View results in GitHub Actions tab
- Extend with deployment steps

## Usage Guide

### *Initial Setup*

#### 1. Prerequisites

```
```bash
# Install Docker Desktop
```

```
# Install Python 3.9
```

```
# Install Git
```

```
...
```

## 2. Clone Repository

```
```bash
```

```
git clone <repository-url>
```

```
cd bike-sharing-mlops
```

```
...
```

## 3. Prepare Data

```
```bash
```

```
# Download data to data/bike_sharing_raw.csv
```

```
# Or use DVC: dvc pull
```

```
...
```

## 4. Start Services

```
```bash
```

```
docker-compose up -d
```

```
...
```

## 5. Access Services

- Airflow UI: http://localhost:8081 (admin/admin)

- MLflow UI: http://localhost:5000

- Streamlit App: http://localhost:8501

- Prometheus: http://localhost:9090

- Grafana: http://localhost:3000

## 6. Start API on Windows

```
```bash
```

```
python src/api.py
```

```
...
```

## *Running the Pipeline*

### *Option 1: Airflow (Automated)*

1. Open Airflow UI: http://localhost:8081
2. Enable DAG: bike\_sharing\_final\_pipeline\_v4
3. Click "Trigger DAG" button
4. Monitor task progress in Graph view

### *Option 2: Manual Execution*

```
# Step 1: Ingest data
python src/ingestion.py

# Step 2: Preprocess
python src/preprocessing.py

# Step 3: Train model
python src/train.py

# Step 4: Evaluate
python src/evaluate.py

# Step 5: Make predictions
python src/api.py # In separate terminal
```

## ***Making Predictions***

### ***Via Streamlit UI***

1. Open <http://localhost:8501>
2. Adjust sliders for features
3. Click "Predict Bike Demand"
4. View result

### ***Via API (Browser)***

```
http://localhost:9999/predict?temp=0.8&hr=18&season=2
```

### ***Via API (Python)***

```
import requests

params = {
    "season": 2,
    "mnth": 6,
    "hr": 18,
    "temp": 0.8,
    "hum": 0.6,
    "windspeed": 0.2
}

response = requests.get("http://localhost:9999/predict", params=params)
print(response.json())
```

### ***Via API (cURL)***

```
curl "http://localhost:9999/predict?temp=0.8&hr=18"
```

## ***Monitoring***

### ***View Drift Reports***

1. Run monitoring DAG in Airflow
2. Open Streamlit app
3. Click "View Data Drift Report"
4. Review Evidently AI report

### ***View Prometheus Metrics***

1. Open <http://localhost:9090>
2. Query: `http_requests_total`
3. View request counts and latencies

### ***View Grafana Dashboards***

1. Open <http://localhost:3000>
2. Login: admin/admin
3. Add Prometheus data source
4. Create dashboards

### ***Troubleshooting***

#### ***API Connection Failed***

```
# Check if API is running
curl http://localhost:9999/

# Check Docker network
docker network inspect bike-mlops-net

# Restart API
python src/api.py
```

#### ***Airflow Task Failed***

```
# View logs
docker logs airflow_scheduler

# Check task logs in Airflow UI
# Click on failed task > Log
```

#### ***Model Not Found***

```
# Train model manually
python src/train.py

# Check model file
ls models/bike_model.pkl
```

#### ***Database Connection Error***

```
# Check PostgreSQL
docker ps | grep postgres

# Restart database
docker-compose restart postgres
```

## **Deployment Instructions**

### ***Local Development***

- Use docker-compose for all services
- Run API on Windows host

- Access via localhost

## *Production Deployment*

### *AWS Deployment*

#### **1. Setup ECS/EKS**

```
```bash
# Push images to ECR
docker tag bike-sharing-mlops:latest <ecr-url>
docker push <ecr-url>
````
```

#### **2. Configure RDS**

- Replace PostgreSQL with RDS
- Update connection strings

#### **3. Setup S3**

- Replace LocalStack with real S3
- Update bucket configurations

#### **4. Configure ALB**

- Load balancer for API
- SSL certificates

### *Azure Deployment*

#### **1. Azure Container Instances**

```
```bash
az container create --resource-group mygroup --name bike-mlops --image <acr-url>
````
```

#### **2. Azure Database for PostgreSQL**

- Managed database service
- Update connection strings

#### **3. Azure Blob Storage**

- Replace LocalStack
- Update storage configs

### *Kubernetes Deployment*

#### **1. Create Kubernetes manifests**

```
```yaml
# deployment.yaml
```

```
# service.yaml
```

```
# ingress.yaml
```

```
...
```

## 2. Deploy

```
```bash
```

```
kubectl apply -f k8s/
```

```
...
```

## 3. Configure Helm charts

- Package application
- Manage releases

## *Environment Variables*

```
# Production .env file
DATABASE_URL=postgresql://user:pass@prod-db:5432/mlops
MLFLOW_TRACKING_URI=http://mlflow.prod:5000
S3_BUCKET=prod-bike-sharing
SLACK_WEBHOOK=https://hooks.slack.com/...
```

## *Security Considerations*

- Use secrets management (AWS Secrets Manager, Azure Key Vault)
- Enable SSL/TLS for all services
- Implement authentication (OAuth, JWT)
- Network isolation (VPC, Security Groups)
- Regular security updates
- Audit logging

## *Scaling Considerations*

- Horizontal scaling for API (multiple replicas)
- Database read replicas
- Caching layer (Redis)
- CDN for static assets
- Auto-scaling based on metrics

## **Key Concepts Explained**

### *MLOps Pipeline*

1. **Data Ingestion:** Load raw data
2. **Data Validation:** Check quality

3. **Preprocessing:** Clean and split
4. **Training:** Build model
5. **Evaluation:** Measure performance
6. **Deployment:** Serve predictions
7. **Monitoring:** Track drift and performance
8. **Retraining:** Update model when needed

## ***Data Drift***

- **What**: Changes in data distribution over time
- **Why Important**: Model trained on old data may not work on new data
- **Detection**: Evidently AI compares current vs reference data
- **Action**: Retrain model when drift detected

## ***Model Registry***

- **What**: Centralized storage for model versions
- **Why Important**: Track which model is in production
- **MLflow**: Provides versioning and aliasing
- **Usage**: Load model by alias ("champion")

## ***Experiment Tracking***

- **What**: Recording model parameters and metrics
- **Why Important**: Compare different models
- **MLflow**: Tracks runs, parameters, metrics, artifacts
- **Usage**: View all experiments in MLflow UI

## ***Containerization***

- **What**: Packaging application with dependencies
- **Why Important**: Consistent environment across machines
- **Docker**: Creates isolated containers
- **Usage**: `docker-compose up` starts all services

## ***Orchestration***

- **What**: Automating workflow execution
- **Why Important**: Scheduled and reliable pipelines
- **Airflow**: DAGs define task dependencies
- **Usage**: Schedule daily retraining

## ***API Serving***

- **What**: Exposing model as web service
- **Why Important**: Applications can get predictions

- **FastAPI**: Modern, fast Python framework
- **Usage**: HTTP GET/POST requests

## ***Monitoring***

- **What**: Tracking system health and performance
- **Why Important**: Detect issues early
- **Prometheus**: Metrics collection
- **Grafana**: Visualization
- **Usage**: Dashboards and alerts

# **Best Practices**

## ***Code Quality***

- Use linting (flake8, pylint)
- Write docstrings
- Type hints for functions
- Unit tests for critical functions
- Code reviews

## ***Data Management***

- Version data with DVC
- Validate data quality
- Document data schema
- Handle missing values
- Monitor data drift

## ***Model Development***

- Track experiments in MLflow
- Use cross-validation
- Feature engineering
- Hyperparameter tuning
- Model versioning

## ***Deployment***

- Containerize applications
- Use environment variables
- Implement health checks
- Graceful error handling

- Logging and monitoring

## ***Security***

- Never commit secrets
- Use environment variables
- Implement authentication
- Regular updates
- Audit logs

## ***Documentation***

- README for setup
- Code comments
- API documentation
- Architecture diagrams
- Runbooks for operations

## **Conclusion**

This project demonstrates a complete MLOps pipeline with:

- Automated training and deployment
- Continuous monitoring and retraining
- Scalable architecture
- Production-ready practices

### **Key Takeaways:**

1. Automation reduces manual errors
2. Monitoring prevents model degradation
3. Containerization ensures consistency
4. Version control tracks changes
5. CI/CD enables rapid iteration

### **Next Steps:**

- Add more features to model
- Implement A/B testing
- Add more monitoring metrics
- Scale to production workload
- Implement model explainability

# Appendix

## Useful Commands

```
# Docker
docker-compose up -d          # Start all services
docker-compose down           # Stop all services
docker-compose logs -f        # View logs
docker ps                      # List containers
docker exec -it <container> bash # Enter container

# Airflow
airflow dags list             # List DAGs
airflow tasks test <dag> <task> <date> # Test task
airflow dags trigger <dag>     # Trigger DAG

# MLflow
mlflow ui                      # Start UI
mlflow models serve -m <model> # Serve model

# DVC
dvc init                         # Initialize DVC
dvc add data/                     # Track data
dvc push                          # Upload data
dvc pull                          # Download data

# Git
git add .                         # Stage changes
git commit -m "message"          # Commit
git push                          # Push to remote
```

## Port Reference

- 5432: PostgreSQL
- 4566: LocalStack (S3)
- 5000: MLflow
- 8081: Airflow
- 8501: Streamlit
- 9999: FastAPI
- 9090: Prometheus
- 3000: Grafana

## File Structure

```
bike-sharing-mlops/
    .dvc/                  # DVC configuration
    .github/workflows/      # CI/CD pipelines
    dags/                  # Airflow DAGs
    data/
        raw/                # Data files (not in Git)
        processed/
    models/                # Model files (not in Git)
    src/
        api.py
        app.py
        train.py
        predict.py
        ...
    docker-compose.yml      # Service orchestration
    Dockerfile              # Streamlit image
    Dockerfile.airflow       # Airflow image
    prometheus.yml          # Monitoring config
    README.md               # Project overview
```

**Document Version:** 1.0

**Last Updated:** February 15, 2026

**Author:** MLOps Team

**Contact:** [Your Contact Information]